

BigDataAnalytics_Project

May 23, 2026

1 Big Data Analytics — Project 1

1.1 Analyzing Sales Data with Dask

Dataset: Sample Sales Data (2,823 rows × 25 columns)

Tools: Dask, Pandas, Scikit-learn, Matplotlib, Seaborn

```
[47]: import dask
      from dask.distributed import Client
      import dask.dataframe as dd
      import pandas as pd
      import numpy as np
      import matplotlib.pyplot as plt
      import seaborn as sns
      import warnings
      warnings.filterwarnings('ignore')

      plt.rcParams['figure.figsize'] = (10, 5)
      sns.set_theme(style='whitegrid')
      print('Libraries loaded ')
```

Libraries loaded

1.2 Task 1 — Data Preparation

1.2.1 1.1 Load dataset using Dask DataFrames

```
[6]: df = pd.read_csv('data/sales_data_sample.csv', encoding='latin1')
      dtypes = df.dtypes.to_dict()

[7]: # Load with Dask for parallel processing
      ddf = dd.read_csv('data/sales_data_sample.csv', encoding='latin1', dtype=dtypes)

      print('Dask DataFrame info:')
      print(f'  Number of partitions: {ddf.npartitions}')
      print(f'  Columns: {list(ddf.columns)}')
      print()
```

```
# Compute shape
n_rows = len(ddf)
n_cols = len(ddf.columns)
print(f' Shape: ({n_rows}, {n_cols})')

# Preview first 5 rows
ddf.head(5)
```

Dask DataFrame info:

Number of partitions: 1

Columns: ['ORDERNUMBER', 'QUANTITYORDERED', 'PRICEEACH', 'ORDERLINENUMBER', 'SALES', 'ORDERDATE', 'STATUS', 'QTR_ID', 'MONTH_ID', 'YEAR_ID', 'PRODUCTLINE', 'MSRP', 'PRODUCTCODE', 'CUSTOMERNAME', 'PHONE', 'ADDRESSLINE1', 'ADDRESSLINE2', 'CITY', 'STATE', 'POSTALCODE', 'COUNTRY', 'TERRITORY', 'CONTACTLASTNAME', 'CONTACTFIRSTNAME', 'DEALSIZE']

Shape: (2823, 25)

```
[7]:  ORDERNUMBER  QUANTITYORDERED  PRICEEACH  ORDERLINENUMBER  SALES  \
0      10107             30      95.70             2  2871.00
1      10121             34      81.35             5  2765.90
2      10134             41      94.74             2  3884.34
3      10145             45      83.26             6  3746.70
4      10159             49     100.00            14  5205.27
```

```
      ORDERDATE  STATUS  QTR_ID  MONTH_ID  YEAR_ID  ...  \
0  2/24/2003 0:00  Shipped      1         2     2003  ...
1  5/7/2003 0:00  Shipped      2         5     2003  ...
2  7/1/2003 0:00  Shipped      3         7     2003  ...
3  8/25/2003 0:00  Shipped      3         8     2003  ...
4 10/10/2003 0:00  Shipped      4        10     2003  ...
```

```
      ADDRESSLINE1  ADDRESSLINE2      CITY  STATE  \
0  897 Long Airport Avenue      <NA>      NYC    NY
1      59 rue de l'Abbaye      <NA>     Reims  <NA>
2  27 rue du Colonel Pierre Avia      <NA>     Paris  <NA>
3      78934 Hillside Dr.      <NA>    Pasadena   CA
4      7734 Strong St.      <NA>  San Francisco   CA
```

```
      POSTALCODE  COUNTRY  TERRITORY  CONTACTLASTNAME  CONTACTFIRSTNAME  DEALSIZE
0      10022      USA      <NA>           Yu           Kwai      Small
1      51100  France      EMEA      Henriot           Paul      Small
2      75508  France      EMEA      Da Cunha      Daniel      Medium
3      90003      USA      <NA>      Young           Julie      Medium
4      <NA>      USA      <NA>      Brown           Julie      Medium
```

[5 rows x 25 columns]

1.2.2 1.2 Data Cleaning — Missing Values, Outliers, Inconsistencies

```
[8]: # Check missing values
missing = ddf.isnull().sum().compute()
missing_pct = (missing / len(ddf) * 100).round(2)


missing_df = pd.DataFrame({
    'Missing Count': missing,
    'Missing %': missing_pct
}).query('`Missing Count` > 0')

print('Columns with missing values:')
print(missing_df)

# Visualize missing values
fig, ax = plt.subplots(figsize=(8, 4))
missing_df['Missing %'].plot(kind='bar', ax=ax, color='salmon')
ax.set_title('Missing Value Percentage by Column')
ax.set_ylabel('Missing %')
ax.set_xlabel('Column')
plt.xticks(rotation=30)
plt.tight_layout()
plt.show()
```

Columns with missing values:

	Missing Count	Missing %
ADDRESSLINE2	2521	89.30
STATE	1486	52.64
POSTALCODE	76	2.69
TERRITORY	1074	38.04



```
[9]: # Handle missing values
# ADDRESSLINE2 → not critical, fill with 'N/A'
# STATE → fill with 'Unknown'
# POSTALCODE → fill with 'Unknown'
# TERRITORY → fill with 'Unknown'

ddf = ddf.fillna({
    'ADDRESSLINE2': 'N/A',
    'STATE': 'Unknown',
    'POSTALCODE': 'Unknown',
    'TERRITORY': 'Unknown'
})

# Convert ORDERDATE to datetime
ddf['ORDERDATE'] = dd.to_datetime(ddf['ORDERDATE'])

# Verify no missing values remain in key columns
remaining_missing = ddf[['SALES', 'QUANTITYORDERED', 'PRICEEACH']].isnull().sum().
    ↪compute()
print('Missing values in key numeric columns after cleaning:')
print(remaining_missing)
print('\nData types after conversion:')
print(ddf.dtypes)
```

Missing values in key numeric columns after cleaning:

```
SALES          0
QUANTITYORDERED  0
PRICEEACH      0
dtype: int64
```

Data types after conversion:

```
ORDERNUMBER          float64
QUANTITYORDERED      float64
PRICEEACH            float64
ORDERLINENUMBER      float64
SALES                float64
ORDERDATE            datetime64[us]
STATUS               string
QTR_ID              float64
MONTH_ID            float64
YEAR_ID             float64
PRODUCTLINE          string
MSRP                float64
PRODUCTCODE          string
CUSTOMERNAME         string
PHONE               string
```

```

ADDRESSLINE1          string
ADDRESSLINE2          string
CITY                   string
STATE                  string
POSTALCODE             string
COUNTRY                string
TERRITORY              string
CONTACTLASTNAME        string
CONTACTFIRSTNAME       string
DEALSIZE               string
dtype: object

```

```

[12]: # Detect Outliers using IQR on SALES and QUANTITYORDERED
df_pd = ddf.compute() # bring to pandas for IQR calculation

outlier_bounds = {}
for col in ['SALES', 'QUANTITYORDERED', 'PRICEEACH']:
    Q1 = df_pd[col].quantile(0.25)
    Q3 = df_pd[col].quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
    outliers = df_pd[(df_pd[col] < lower) | (df_pd[col] > upper)]
    outlier_bounds[col] = (lower, upper)
    print(f'{col}: {len(outliers)} outliers detected '
          f'bounds: [{lower:.2f}, {upper:.2f}]')


# Visualize BEFORE handling
fig, axes = plt.subplots(1, 3, figsize=(14, 4))
for ax, col in zip(axes, ['SALES', 'QUANTITYORDERED', 'PRICEEACH']):
    df_pd.boxplot(column=col, ax=ax)
    ax.set_title(f'{col} Distribution')
plt.suptitle('Outlier Detection - Box Plots (Before Handling)', fontsize=13)
plt.tight_layout()
plt.show()

```

```

SALES: 81 outliers detected (bounds: [-1253.42, 7964.85])
QUANTITYORDERED: 8 outliers detected (bounds: [3.00, 67.00])
PRICEEACH: 0 outliers detected (bounds: [22.15, 146.71])

```



```
[13]: # Handle Outliers
# SALES (81 outliers)      + Winsorize: cap at IQR bounds
# Reason: large orders are real business events, don't delete them but cap to
# → reduce skew effect on modeling
# QUANTITYORDERED (8 rows) + Keep as-is
# Reason: only 8 rows, very likely genuine bulk orders
# PRICEEACH (0 outliers)   + No action needed

lower_sales, upper_sales = outlier_bounds['SALES']

print(f'\nSALES before Winsorize - min: {df_pd["SALES"].min():.2f}, '
      f'max: {df_pd["SALES"].max():.2f}, '
      f'mean: {df_pd["SALES"].mean():.2f}')

# Winsorize: clip values at IQR bounds
df_pd['SALES'] = df_pd['SALES'].clip(lower=lower_sales, upper=upper_sales)

print(f'SALES after Winsorize - min: {df_pd["SALES"].min():.2f}, '
      f'max: {df_pd["SALES"].max():.2f}, '
      f'mean: {df_pd["SALES"].mean():.2f}')

remaining = len(df_pd[
    (df_pd['SALES'] < lower_sales) | (df_pd['SALES'] > upper_sales)
])
print(f'Remaining SALES outliers after Winsorize: {remaining}')

print(f'\nQUANTITYORDERED - 8 outliers kept (genuine bulk orders):')
print(df_pd[df_pd['QUANTITYORDERED'] > outlier_bounds['QUANTITYORDERED'][1]]
      [['ORDERNUMBER', 'CUSTOMERNAME', 'QUANTITYORDERED', 'SALES']])

print(f'\nPRICEEACH - 0 outliers, no action needed.')

# Sync cleaned pandas DataFrame back to Dask
ddf = dd.from_pandas(df_pd, npartitions=4)
```

```

print(f'\nDask DataFrame updated with cleaned data. Partitions: {ddf.
↳npartitions}')

# Visualize AFTER handling
fig, axes = plt.subplots(1, 3, figsize=(14, 4))
for ax, col in zip(axes, ['SALES', 'QUANTITYORDERED', 'PRICEEACH']):
    df_pd.boxplot(column=col, ax=ax)
    ax.set_title(f'{col} Distribution')
plt.suptitle('Outlier Detection - Box Plots (After Handling)', fontsize=13)
plt.tight_layout()
plt.show()

```

SALES before Winsorize - min: 482.13, max: 14082.80, mean: 3553.89

SALES after Winsorize - min: 482.13, max: 7964.85, mean: 3519.25

Remaining SALES outliers after Winsorize: 0

QUANTITYORDERED - 8 outliers kept (genuine bulk orders):

	ORDERNUMBER	CUSTOMERNAME	QUANTITYORDERED	SALES
418	10405	Mini Caravy	97	7964.855
598	10407	The Sharp Gifts Warehouse	76	7964.855
1666	10412	Euro Shopping Channel	70	7964.855
1714	10407	The Sharp Gifts Warehouse	76	7182.000
1995	10405	Mini Caravy	76	7964.855
1996	10419	Salzburg Collectables	70	7964.855
2586	10401	Tekni Collectables Inc.	85	7543.750
2689	10401	Tekni Collectables Inc.	77	7084.000

PRICEEACH - 0 outliers, no action needed.

Dask DataFrame updated with cleaned data. Partitions: 4



```

[14]: # Check for inconsistencies
print('Unique STATUS values:')
print(df_pd['STATUS'].value_counts())

```

```

print('\nUnique DEALSIZE values:')
print(df_pd['DEALSIZE'].value_counts())

print('\nUnique PRODUCTLINE values:')
print(df_pd['PRODUCTLINE'].value_counts())

print('\nYear range in data:')
print(df_pd['YEAR_ID'].value_counts().sort_index())

```

Unique STATUS values:

STATUS

Shipped	2617
---------	------

Cancelled	60
-----------	----

Resolved	47
----------	----

On Hold	44
---------	----

In Process	41
------------	----

Disputed	14
----------	----

Name: count, dtype: int64[pyarrow]

Unique DEALSIZE values:

DEALSIZE

Medium	1384
--------	------

Small	1282
-------	------

Large	157
-------	-----

Name: count, dtype: int64[pyarrow]

Unique PRODUCTLINE values:

PRODUCTLINE

Classic Cars	967
--------------	-----

Vintage Cars	607
--------------	-----

Motorcycles	331
-------------	-----

Planes	306
--------	-----

Trucks and Buses	301
------------------	-----

Ships	234
-------	-----

Trains	77
--------	----

Name: count, dtype: int64[pyarrow]

Year range in data:

YEAR_ID

2003	1000
------	------

2004	1345
------	------

2005	478
------	-----

Name: count, dtype: int64

1.3 Task 2 — Exploratory Data Analysis (EDA)

1.3.1 2.1 Summary Statistics

```
[15]: # Summary statistics using Dask
print('=== Summary Statistics (Dask) ===')
summary = ddf[['SALES', 'QUANTITYORDERED', 'PRICEEACH', 'MSRP']].describe().
    ↪compute()
print(summary.round(2))
```

```
=== Summary Statistics (Dask) ===
```

	SALES	QUANTITYORDERED	PRICEEACH	MSRP
count	2823.00	2823.00	2823.00	2823.00
mean	3519.25	35.09	83.66	100.72
std	1731.68	9.74	20.17	40.19
min	482.13	6.00	26.88	33.00
25%	2612.48	28.00	78.89	80.00
50%	3609.60	36.00	100.00	108.50
75%	5904.39	43.00	100.00	140.00
max	7964.86	97.00	100.00	214.00

1.3.2 2.2 Sales Distribution — Products, Customers, Time

```
[16]: # Sales by Product Line
sales_by_product = ddf.groupby('PRODUCTLINE')['SALES'].sum().compute().
    ↪sort_values(ascending=False)


fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Bar chart
sales_by_product.plot(kind='bar', ax=axes[0], color='steelblue')
axes[0].set_title('Total Sales by Product Line')
axes[0].set_ylabel('Total Sales ($)')
axes[0].set_xlabel('Product Line')
axes[0].tick_params(axis='x', rotation=30)

# Pie chart
axes[1].pie(sales_by_product, labels=sales_by_product.index,
            autopct='%1.1f%%', startangle=140)
axes[1].set_title('Sales Share by Product Line')

plt.suptitle('Sales Distribution by Product Line', fontsize=13)
plt.tight_layout()
plt.show()

print('Sales by Product Line:')
print(sales_by_product)
```



Sales by Product Line:

PRODUCTLINE


Classic Cars	3865484.790
Vintage Cars	1879004.410
Motorcycles	1154421.530
Trucks and Buses	1125623.055
Planes	970631.635
Ships	714437.130
Trains	225231.275

Name: SALES, dtype: float64

```
[17]: # Top 10 Customers by Total Sales
top_customers = (
    ddf.groupby('CUSTOMERNAME')['SALES']
        .sum()
        .compute()
        .sort_values(ascending=False)
        .head(10)
)

fig, ax = plt.subplots(figsize=(11, 5))
top_customers.plot(kind='barh', ax=ax, color='darkorange')
ax.set_title('Top 10 Customers by Total Sales')
ax.set_xlabel('Total Sales ($)')
ax.invert_yaxis()
plt.tight_layout()
plt.show()

print('Top 10 Customers:')
print(top_customers.reset_index().rename(columns={'SALES': 'Total Sales ($)'}))
```



Top 10 Customers:

	CUSTOMERNAME	Total Sales (\$)
0	Euro Shopping Channel	902217.880
1	Mini Gifts Distributors Ltd.	647302.655
2	Australian Collectors, Co.	197886.230
3	Muscle Machine Inc	194248.690
4	La Rochelle Gifts	178050.155
5	Dragon Souveniers, Ltd.	168541.535
6	Land of Toys Inc.	164069.440
7	AV Stores, Co.	157807.810
8	Anna's Decorations, Ltd	153110.990
9	The Sharp Gifts Warehouse	152000.295

1.3.3 2.3 Sales Trends Over Time

```
[18]: # Monthly Sales Trend
monthly_sales = (
    ddf.groupby(['YEAR_ID', 'MONTH_ID'])['SALES']
        .sum()
        .compute()
        .reset_index()
        .sort_values(['YEAR_ID', 'MONTH_ID'])
)
monthly_sales['YearMonth'] = (monthly_sales['YEAR_ID'].astype(str) + '-'
                             + monthly_sales['MONTH_ID'].astype(str).str.
                             zfill(2))


fig, ax = plt.subplots(figsize=(14, 5))
for year in sorted(monthly_sales['YEAR_ID'].unique()):
    subset = monthly_sales[monthly_sales['YEAR_ID'] == year]
```

```

ax.plot(subset['MONTH_ID'], subset['SALES'], marker='o', label=str(year))

ax.set_title('Monthly Sales Trend by Year')
ax.set_xlabel('Month')
ax.set_ylabel('Total Sales ($)')
ax.set_xticks(range(1,13))
ax.set_xticklabels(['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun',
                    'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec'])
ax.legend(title='Year')
plt.tight_layout()
plt.show()

```




```

[19]: # Sales by Country
sales_by_country = (
    ddf.groupby('COUNTRY')['SALES']
        .sum()
        .compute()
        .sort_values(ascending=False)
        .head(10)
)

fig, ax = plt.subplots(figsize=(10, 5))
sales_by_country.plot(kind='bar', ax=ax, color='seagreen')
ax.set_title('Top 10 Countries by Total Sales')
ax.set_ylabel('Total Sales ($)')
ax.set_xlabel('Country')
ax.tick_params(axis='x', rotation=30)
plt.tight_layout()
plt.show()

```




```
[20]: # Deal Size Distribution
deal_sales = ddf.groupby('DEALSIZE')['SALES'].sum().compute()

fig, axes = plt.subplots(1, 2, figsize=(12, 5))
deal_sales.plot(kind='bar', ax=axes[0],
                color=['steelblue', 'darkorange', 'firebrick'])
axes[0].set_title('Total Sales by Deal Size')
axes[0].set_ylabel('Total Sales ($)')
axes[0].tick_params(axis='x', rotation=0)


deal_count = ddf.groupby('DEALSIZE')['ORDERNUMBER'].count().compute()
deal_count.plot(kind='pie', ax=axes[1], autopct='%1.1f%%', startangle=140)
axes[1].set_title('Order Count by Deal Size')
axes[1].set_ylabel('')

plt.suptitle('Deal Size Analysis', fontsize=13)
plt.tight_layout()
plt.show()
```



1.3.4 2.4 Correlation heatmap

```
[55]: numeric_cols = ['SALES', 'QUANTITYORDERED', 'PRICEEACH', 'MSRP']
sns.heatmap(df_pd[numeric_cols].corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Feature Correlation Heatmap')
plt.show()
```



1.4 Task 3 — Feature Engineering

```
[21]: # Feature 1: Total spending per customer
customer_spending = (
    ddf.groupby('CUSTOMERNAME')['SALES']
    .sum()
    .compute()
    .reset_index()
    .rename(columns={'SALES': 'TOTAL_SPENDING'})
)
print('Total spending per customer (top 5):')
print(customer_spending.sort_values('TOTAL_SPENDING', ascending=False).head())

# Feature 2: Average sales per product line
avg_sales_product = (
```

```

ddf.groupby('PRODUCTLINE')['SALES']
    .mean()
    .compute()
    .reset_index()
    .rename(columns={'SALES': 'AVG_SALES'})
)
print('\nAverage sales per product line:')
print(avg_sales_product.sort_values('AVG_SALES', ascending=False))

```

Total spending per customer (top 5):

	CUSTOMERNAME	TOTAL_SPENDING
32	Euro Shopping Channel	902217.880
54	Mini Gifts Distributors Ltd.	647302.655
6	Australian Collectors, Co.	197886.230
57	Muscle Machine Inc	194248.690
43	La Rochelle Gifts	178050.155

Average sales per product line:

	PRODUCTLINE	AVG_SALES
0	Classic Cars	3997.398956
3	Trucks and Buses	3739.611478
1	Motorcycles	3487.678338
2	Planes	3171.998807
4	Vintage Cars	3095.559160
5	Ships	3053.150128
6	Trains	2925.081494

```

[22]: # Feature 3: Time-based features (Dask parallel computation)
# Add quarter, day of week, is_weekend
ddf['QUARTER'] = ddf['ORDERDATE'].dt.quarter
ddf['DAY_OF_WEEK'] = ddf['ORDERDATE'].dt.dayofweek # 0=Mon, 6=Sun
ddf['IS_WEEKEND'] = (ddf['DAY_OF_WEEK'] >= 5).astype(int)

# Feature 4: Revenue per unit
ddf['REVENUE_PER_UNIT'] = ddf['SALES'] / ddf['QUANTITYORDERED']

# Feature 5: Price discount ratio (MSRP vs actual price)
ddf['DISCOUNT_RATIO'] = ((ddf['MSRP'] - ddf['PRICEEACH']) / ddf['MSRP']).
    .clip(0, 1)

print('New features added:')
print(ddf[['ORDERDATE', 'QUARTER', 'DAY_OF_WEEK', 'IS_WEEKEND',
    'REVENUE_PER_UNIT', 'DISCOUNT_RATIO']].head(5))

```

New features added:

	ORDERDATE	QUARTER	DAY_OF_WEEK	IS_WEEKEND	REVENUE_PER_UNIT	\
0	2003-02-24	1	0	0	95.70	
1	2003-05-07	2	2	0	81.35	


2	2003-07-01	3	1	0	94.74
3	2003-08-25	3	0	0	83.26
4	2003-10-10	4	4	0	106.23

	DISCOUNT_RATIO
0	0.000000
1	0.143684
2	0.002737
3	0.123579
4	0.000000

```
[23]: # Seasonality analysis using new QUARTER feature
quarterly_sales = (
    ddf.groupby(['YEAR_ID', 'QUARTER'])['SALES']
        .sum()
        .compute()
        .reset_index()
)

fig, ax = plt.subplots(figsize=(10, 5))
for year in sorted(quarterly_sales['YEAR_ID'].unique()):
    subset = quarterly_sales[quarterly_sales['YEAR_ID'] == year]
    ax.plot(subset['QUARTER'], subset['SALES'], marker='s', label=str(year))

ax.set_title('Quarterly Sales by Year (Seasonality)')
ax.set_xlabel('Quarter')
ax.set_ylabel('Total Sales ($)')
ax.set_xticks([1, 2, 3, 4])
ax.set_xticklabels(['Q1', 'Q2', 'Q3', 'Q4'])
ax.legend(title='Year')
plt.tight_layout()
plt.show()
```



1.5 Task 4 — Advanced Analytics

1.5.1 4.1 Customer Segmentation with K-Means Clustering

```
[29]: from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score

# Build RFM-style features per customer
df_pd = ddf.compute()

rfm = df_pd.groupby('CUSTOMERNAME').agg(
    TOTAL_SALES = ('SALES', 'sum'),
    ORDER_COUNT = ('ORDERNUMBER', 'nunique'),
    AVG_ORDER_VAL = ('SALES', 'mean')
).reset_index()

# Scale features
scaler = StandardScaler()
X_rfm = scaler.
    ↪ fit_transform(rfm[['TOTAL_SALES', 'ORDER_COUNT', 'AVG_ORDER_VAL']])


# Elbow method to find optimal k
inertias = []
K_range = range(2, 9)
for k in K_range:
    km = KMeans(n_clusters=k, random_state=42, n_init=10)
```

```

km.fit(X_rfm)
inertias.append(km.inertia_)

fig, ax = plt.subplots(figsize=(8, 4))
ax.plot(list(K_range), inertias, marker='o', color='steelblue')
ax.set_title('Elbow Method - Optimal Number of Clusters')
ax.set_xlabel('Number of Clusters (k)')
ax.set_ylabel('Inertia')
plt.tight_layout()
plt.show()

```



```

[30]: from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
import matplotlib.pyplot as plt
import numpy as np
import matplotlib.cm as cm

# Step 1: Check outliers in RFM before clustering
print('RFM stats before outlier removal:')
print(rfm[['TOTAL_SALES', 'ORDER_COUNT', 'AVG_ORDER_VAL']].describe().round(2))

# Step 2: Remove extreme outliers using IQR on ORDER_COUNT
Q1 = rfm['ORDER_COUNT'].quantile(0.25)
Q3 = rfm['ORDER_COUNT'].quantile(0.75)
IQR = Q3 - Q1
upper_bound = Q3 + 1.5 * IQR

```

```

# Separate outlier customers (keep them labeled separately)
rfm_outliers = rfm[rfm['ORDER_COUNT'] > upper_bound].copy()
rfm_clean     = rfm[rfm['ORDER_COUNT'] <= upper_bound].copy()

print(f'\nIQR upper bound for ORDER_COUNT: {upper_bound:.1f}')
print(f'Extreme customers removed from clustering: {len(rfm_outliers)}')
print(rfm_outliers[['CUSTOMERNAME', 'TOTAL_SALES', 'ORDER_COUNT', 'AVG_ORDER_VAL']])
print(f'Customers used for clustering: {len(rfm_clean)}')

# Step 3: Scale features
scaler = StandardScaler()
X_clean = scaler.fit_transform(
    rfm_clean[['TOTAL_SALES', 'ORDER_COUNT', 'AVG_ORDER_VAL']]
)

# Step 4: Elbow + Silhouette to find optimal k
inertias = []
sil_scores = []
K_range = range(2, 8)

for k in K_range:
    km = KMeans(n_clusters=k, random_state=42, n_init=10)
    lbl = km.fit_predict(X_clean)
    inertias.append(km.inertia_)
    sil_scores.append(silhouette_score(X_clean, lbl))

fig, axes = plt.subplots(1, 2, figsize=(13, 4))

axes[0].plot(list(K_range), inertias, marker='o', color='steelblue')
axes[0].set_title('Elbow Method')
axes[0].set_xlabel('Number of Clusters (k)')
axes[0].set_ylabel('Inertia')

axes[1].plot(list(K_range), sil_scores, marker='s', color='darkorange')
axes[1].set_title('Silhouette Score by k')
axes[1].set_xlabel('Number of Clusters (k)')
axes[1].set_ylabel('Silhouette Score')

plt.suptitle('Finding Optimal k - Elbow & Silhouette', fontsize=13)
plt.tight_layout()
plt.show()

# Print scores to help pick k
print('\nSilhouette Scores:')
for k, s in zip(K_range, sil_scores):
    print(f' k={k}: {s:.4f}')

```

```

# Step 5: Apply K-Means with best k
# Pick k with highest silhouette score
k_optimal = list(K_range)[np.argmax(sil_scores)]
print(f'\nBest k by silhouette score: {k_optimal}')

kmeans = KMeans(n_clusters=k_optimal, random_state=42, n_init=10)
rfm_clean['CLUSTER'] = kmeans.fit_predict(X_clean)

# Label outlier customers as separate group
rfm_outliers['CLUSTER'] = -1 # -1 = VIP/Outlier group

# Combine back
rfm_final = pd.concat([rfm_clean, rfm_outliers], ignore_index=True)

# Step 6: Cluster profiles
sil_final = silhouette_score(X_clean, rfm_clean['CLUSTER'])
print(f'Silhouette Score (k={k_optimal}): {sil_final:.4f}')

print('\nCluster Profiles (main clusters):')
profile = rfm_clean.groupby('CLUSTER')[
    ['TOTAL_SALES', 'ORDER_COUNT', 'AVG_ORDER_VAL']
].mean().round(2)
print(profile)

# Add descriptive labels based on profile
cluster_labels = {}
for c in range(k_optimal):
    row = profile.loc[c]
    if row['TOTAL_SALES'] == profile['TOTAL_SALES'].max():
        cluster_labels[c] = f'Cluster {c} - High Value'
    elif row['TOTAL_SALES'] == profile['TOTAL_SALES'].min():
        cluster_labels[c] = f'Cluster {c} - Low Value'
    else:
        cluster_labels[c] = f'Cluster {c} - Mid Value'

print('\nCluster Labels:')
for c, label in cluster_labels.items():
    print(f' {label}')
print(' Cluster -1 - VIP / Extreme customers (excluded from K-Means)')

# Step 7: Final scatter plot
# Generate enough colors dynamically based on k_optimal
colors = [cm.tab10(i) for i in range(k_optimal)] # supports up to 10 clusters

fig, ax = plt.subplots(figsize=(10, 6))

# Plot main clusters

```

```

for c in range(k_optimal):
    sub = rfm_clean[rfm_clean['CLUSTER'] == c]
    ax.scatter(sub['ORDER_COUNT'], sub['TOTAL_SALES'],
               label=cluster_labels[c],
               color=colors[c], alpha=0.7, s=60)

# Plot VIP outliers separately
ax.scatter(rfm_outliers['ORDER_COUNT'], rfm_outliers['TOTAL_SALES'],
           label='VIP / Extreme Customers',
           color='darkorange', marker='*', s=300, zorder=5)

# Annotate VIP customers
for _, row in rfm_outliers.iterrows():
    ax.annotate(row['CUSTOMERNAME'],
                (row['ORDER_COUNT'], row['TOTAL_SALES']),
                textcoords='offset points', xytext=(6, 4), fontsize=8)

ax.set_xlabel('Order Count')
ax.set_ylabel('Total Sales ($)')
ax.set_title(f'Customer Segmentation - K-Means (k={k_optimal}) + VIP Detection')
ax.legend(fontsize=9)
plt.tight_layout()
plt.show()

```

RFM stats before outlier removal:

	TOTAL_SALES	ORDER_COUNT	AVG_ORDER_VAL
count	92.00	92.00	92.00
mean	107987.32	3.34	3519.74
std	108991.58	2.92	351.76
min	9129.35	1.00	2499.57
25%	69622.36	2.00	3307.68
50%	86522.61	3.00	3508.98
75%	118088.60	3.00	3750.68
max	902217.88	26.00	4620.54


IQR upper bound for ORDER_COUNT: 4.5

Extreme customers removed from clustering: 6

	CUSTOMERNAME	TOTAL_SALES	ORDER_COUNT	AVG_ORDER_VAL
6	Australian Collectors, Co.	197886.230	5	3597.931455
27	Danish Wholesale Imports	141402.215	5	3927.839306
31	Dragon Souveniers, Ltd.	168541.535	5	3919.570581
33	Euro Shopping Channel	902217.880	26	3483.466718
55	Mini Gifts Distributors Ltd.	647302.655	17	3596.125861
66	Reims Collectables	134723.795	5	3285.946220

Customers used for clustering: 86

Finding Optimal k — Elbow & Silhouette



Silhouette Scores:

k=2: 0.3190
k=3: 0.3101
k=4: 0.3303
k=5: 0.3200
k=6: 0.3377
k=7: 0.3218


Best k by silhouette score: 6
Silhouette Score (k=6): 0.3377

Cluster Profiles (main clusters):

	TOTAL_SALES	ORDER_COUNT	AVG_ORDER_VAL
CLUSTER			
0	94778.69	2.00	3901.10
1	74967.00	3.00	3262.96
2	145511.99	3.71	3480.65
3	46603.60	1.80	2884.82
4	60140.86	2.00	3444.99
5	92542.62	3.20	3863.93

Cluster Labels:

Cluster 0 - Mid Value
Cluster 1 - Mid Value
Cluster 2 - High Value
Cluster 3 - Low Value
Cluster 4 - Mid Value
Cluster 5 - Mid Value
Cluster -1 - VIP / Extreme customers (excluded from K-Means)



1.5.2 4.2 Sales Forecasting with Random Forest + Gradient Boosting

```
[31]: from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.preprocessing import LabelEncoder

# Build feature matrix for predicting QUANTITYORDERED
# Features: PRICEEACH, MSRP, PRODUCTLINE, DEALSIZE, MONTH_ID, QTR_ID, YEAR_ID
# Target : QUANTITYORDERED

df_ml = df_pd[['PRICEEACH', 'MSRP', 'PRODUCTLINE', 'DEALSIZE',
               'MONTH_ID', 'QTR_ID', 'YEAR_ID', 'DISCOUNT_RATIO',
               'QUANTITYORDERED']].copy()

# Encode categorical columns
le_product = LabelEncoder()
le_dealsize = LabelEncoder()
df_ml['PRODUCTLINE_ENC'] = le_product.fit_transform(df_ml['PRODUCTLINE'])
df_ml['DEALSIZE_ENC'] = le_dealsize.fit_transform(df_ml['DEALSIZE'])

feature_cols = ['PRICEEACH', 'MSRP', 'DISCOUNT_RATIO',
                'MONTH_ID', 'QTR_ID', 'YEAR_ID',
                'PRODUCTLINE_ENC', 'DEALSIZE_ENC']
```



```

X = df_ml[feature_cols]
y = df_ml['QUANTITYORDERED']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
print(f'Train size: {len(X_train)} | Test size: {len(X_test)}')

# Model 1: Random Forest
rf = RandomForestRegressor(n_estimators=100, random_state=42, n_jobs=-1)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)

mae_rf = mean_absolute_error(y_test, y_pred_rf)
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))
r2_rf = r2_score(y_test, y_pred_rf)

print(f'\nRandom Forest - Predict QUANTITYORDERED:')
print(f' MAE = {mae_rf:.4f}')
print(f' RMSE = {rmse_rf:.4f}')
print(f' R² = {r2_rf:.4f}')

# Model 2: Gradient Boosting
gb = GradientBoostingRegressor(n_estimators=100, random_state=42)
gb.fit(X_train, y_train)
y_pred_gb = gb.predict(X_test)

mae_gb = mean_absolute_error(y_test, y_pred_gb)
rmse_gb = np.sqrt(mean_squared_error(y_test, y_pred_gb))
r2_gb = r2_score(y_test, y_pred_gb)

print(f'\nGradient Boosting - Predict QUANTITYORDERED:')
print(f' MAE = {mae_gb:.4f}')
print(f' RMSE = {rmse_gb:.4f}')
print(f' R² = {r2_gb:.4f}')

# Model Comparison Table
comparison = pd.DataFrame({
    'Model': ['Random Forest', 'Gradient Boosting'],
    'MAE': [mae_rf, mae_gb],
    'RMSE': [rmse_rf, rmse_gb],
    'R²': [r2_rf, r2_gb]
})
print('\nModel Comparison:')
print(comparison.to_string(index=False))

# Feature Importance (Random Forest)

```

```

importances = pd.Series(rf.feature_importances_, index=feature_cols).
    ↪sort_values(ascending=True)

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Feature importance
importances.plot(kind='barh', ax=axes[0], color='steelblue')
axes[0].set_title('Feature Importance - Random Forest')
axes[0].set_xlabel('Importance Score')

# Actual vs Predicted
axes[1].scatter(y_test, y_pred_rf, alpha=0.4, color='darkorange', s=20)
axes[1].plot([y_test.min(), y_test.max()],
             [y_test.min(), y_test.max()],
             'r--', linewidth=2, label='Perfect Prediction')
axes[1].set_title('Actual vs Predicted - Random Forest')
axes[1].set_xlabel('Actual QUANTITYORDERED')
axes[1].set_ylabel('Predicted QUANTITYORDERED')
axes[1].legend()

plt.suptitle('ML Model - Predicting Order Quantities', fontsize=13)
plt.tight_layout()
plt.show()

```

Train size: 2258 | Test size: 565

Random Forest - Predict QUANTITYORDERED:

```

MAE  = 6.1075
RMSE = 7.7253
R2  = 0.3766

```

Gradient Boosting - Predict QUANTITYORDERED:


```

MAE  = 5.8193
RMSE = 7.1540
R2  = 0.4654

```

Model Comparison:

	Model	MAE	RMSE	R ²
	Random Forest	6.107497	7.725323	0.376598
	Gradient Boosting	5.819277	7.154049	0.465388



1.5.3 4.3 Association Rule Mining — Products Bought Together

```
[32]: from mlxtend.frequent_patterns import apriori, association_rules
      from mlxtend.preprocessing import TransactionEncoder

      # Build basket: each order → list of product lines
      basket = (
          df_pd.groupby('ORDERNUMBER')['PRODUCTLINE']
              .apply(list)
              .reset_index()
      )

      te = TransactionEncoder()
      te_array = te.fit_transform(basket['PRODUCTLINE'])
      basket_df = pd.DataFrame(te_array, columns=te.columns_)

      # Frequent itemsets
      frequent_itemsets = apriori(basket_df, min_support=0.05, use_colnames=True)
      rules = association_rules(frequent_itemsets, metric='lift', min_threshold=1.0)
      rules = rules.sort_values('lift', ascending=False)

      print(f'Frequent itemsets found: {len(frequent_itemsets)}')
      print(f'Association rules found: {len(rules)}')
      print('\nTop 10 rules by lift:')
      print(rules[['antecedents', 'consequents', 'support', 'confidence', 'lift']]
              .head(10).to_string(index=False))
```

Frequent itemsets found: 31

Association rules found: 74

Top 10 rules by lift:

	antecedents			lift
consequents	support	confidence		

```

frozenset({Vintage Cars, Trains})
frozenset({Classic Cars, Ships}) 0.058632    0.692308 5.903846
    frozenset({Classic Cars, Ships})    frozenset({Vintage
Cars, Trains}) 0.058632    0.500000 5.903846
    frozenset({Motorcycles, Vintage Cars})
frozenset({Planes}) 0.058632    0.857143 4.460048
    frozenset({Planes})    frozenset({Motorcycles,
Vintage Cars}) 0.058632    0.305085 4.460048
    frozenset({Ships})    frozenset({Vintage
Cars, Trains}) 0.074919    0.353846 4.178107
    frozenset({Vintage Cars, Trains})
frozenset({Ships}) 0.074919    0.884615 4.178107
frozenset({Classic Cars, Vintage Cars, Trains})
frozenset({Ships}) 0.058632    0.857143 4.048352
    frozenset({Ships}) frozenset({Classic Cars, Vintage
Cars, Trains}) 0.058632    0.276923 4.048352
    frozenset({Classic Cars, Ships})
frozenset({Trains}) 0.058632    0.500000 3.411111
    frozenset({Trains})
frozenset({Classic Cars, Ships}) 0.058632    0.400000 3.411111

```

1.6 Task 5 — Performance Optimization with Dask

```

[33]: import time

# Repartition for optimal parallel processing
print(f'Original partitions: {ddf.npartitions}')

# Repartition to 4 partitions (suitable for a 4-core machine)
ddf_optimized = ddf.repartition(npartitions=4)
print(f'After repartition: {ddf_optimized.npartitions}')

```

Original partitions: 4

After repartition: 4

```

[42]: dtype_spec = {
    'ADDRESSLINE2': 'str',
    'STATE':        'str',
    'POSTALCODE':   'str',
    'TERRITORY':    'str'
}

df_small  = pd.read_csv('data/sales_data_sample.csv', encoding='latin1')
df_large  = pd.concat([df_small] * 500, ignore_index=True)
LARGE_PATH = 'data/sales_data_large.csv'
df_large.to_csv(LARGE_PATH, index=False)

```

```
print(f'Large dataset created: {len(df_large):,} rows × {len(df_large.columns):,} columns')
print(f'Saved to: {LARGE_PATH}')
```

Large dataset created: 1,411,500 rows × 25 columns
 Saved to: data/sales_data_large.csv

```
[43]: #
# 5.1 - Partition Size Tuning
#
print('\n 5.1 Partition Size Tuning  ')

# Load large dataset
ddf_base = dd.read_csv(LARGE_PATH, dtype=dtype_spec)
print(f'Default partitions : {ddf_base.npartitions}')
print(f'Estimated size      : {ddf_base.memory_usage(deep=True).sum().compute() /
      ↪ 1e6:.2f} MB')

# Test different partition sizes and measure performance
partition_tests = [2, 4, 8, 16]
partition_times = []

for n in partition_tests:
    ddf_p = ddf_base.repartition(npartitions=n)
    t0     = time.time()
    _      = ddf_p.groupby('PRODUCTLINE')['SALES'].sum().compute()
    elapsed = time.time() - t0
    partition_times.append(elapsed)
    print(f'  npartitions={n:>2} → {elapsed:.4f}s')

# Find best partition size
best_n = partition_tests[partition_times.index(min(partition_times))]
print(f'\nBest partition size: {best_n} (fastest: {min(partition_times):.4f}s')

# Apply optimal partition
ddf_optimized = ddf_base.repartition(npartitions=best_n)
print(f'ddf_optimized set to {best_n} partitions')

# Plot partition tuning results
fig, ax = plt.subplots(figsize=(8, 4))
bars = ax.bar([str(n) for n in partition_tests], partition_times,
              color=['darkorange' if n == best_n else 'steelblue'
                    for n in partition_tests])
ax.set_title('Execution Time by Number of Partitions')
ax.set_xlabel('Number of Partitions')
ax.set_ylabel('Time (seconds)')
```

```

for bar, val in zip(bars, partition_times):
    ax.text(bar.get_x() + bar.get_width()/2,
            bar.get_height() + 0.001,
            f'{val:.4f}s', ha='center', va='bottom', fontsize=10)
plt.tight_layout()
plt.show()

```

5.1 Partition Size Tuning

Default partitions : 4

Estimated size : 470.05 MB

npartitions= 2 → 6.3913s


npartitions= 4 → 5.7447s

npartitions= 8 → 4.9955s

npartitions=16 → 6.8937s

Best partition size: 8 (fastest: 4.9955s)

ddf_optimized set to 8 partitions



```

[44]: #
# 5.2 - Lazy Evaluation & Efficient Memory Usage
#
print('\n 5.2 Lazy Evaluation & Memory Usage ')

# Lazy evaluation: chain multiple operations WITHOUT computing
lazy_pipeline = (
    ddf_optimized
    .assign(REVENUE_PER_UNIT = ddf_optimized['SALES'] /
    ↪ddf_optimized['QUANTITYORDERED'])

```

```

        .query("STATUS == 'Shipped'")
        .groupby(['PRODUCTLINE', 'YEAR_ID'])['REVENUE_PER_UNIT']
        .mean()
    )

print('Lazy pipeline built - nothing computed yet')
print(f'Type: {type(lazy_pipeline)}')

# Visualize the task graph
try:
    lazy_pipeline.visualize(filename='task_graph.png', rankdir='LR')
    print('Task graph saved → task_graph.png')
    from IPython.display import Image
    display(Image('task_graph.png'))
except Exception:
    print('(graphviz not available - skipping task graph)')

# Trigger computation once
t0 = time.time()
result = lazy_pipeline.compute()
elapsed = time.time() - t0
print(f'\nFull pipeline computed in {elapsed:.4f}s')
print(result.head(10))

# Memory usage: compare eager vs lazy loading
import psutil, os

process = psutil.Process(os.getpid())

# Eager (Pandas): loads everything into RAM immediately
mem_before = process.memory_info().rss / 1e6
df_eager = pd.read_csv(LARGE_PATH)
mem_after = process.memory_info().rss / 1e6
pandas_mem = mem_after - mem_before
print(f'\nMemory Usage Comparison:')
print(f' Pandas (eager load) : +{pandas_mem:.1f} MB loaded into RAM,
↳ immediately')

# Lazy (Dask): only metadata loaded
mem_before = process.memory_info().rss / 1e6
ddf_lazy = dd.read_csv(LARGE_PATH, dtype=dtype_spec)
mem_after = process.memory_info().rss / 1e6
dask_mem = mem_after - mem_before
print(f' Dask (lazy load) : +{dask_mem:.1f} MB (only metadata, data loaded
↳ on demand)')
print(f' Memory saved : {pandas_mem - dask_mem:.1f} MB')

```

```

# Bar chart
fig, ax = plt.subplots(figsize=(7, 4))
ax.bar(['Pandas\n(Eager)', 'Dask\n(Lazy)'],
       [pandas_mem, dask_mem],
       color=['steelblue', 'darkorange'], width=0.4)
ax.set_title('Memory Usage - Eager vs Lazy Loading')
ax.set_ylabel('Memory Increase (MB)')
for i, val in enumerate([pandas_mem, dask_mem]):
    ax.text(i, val + 0.5, f'{val:.1f} MB', ha='center', fontsize=11)
plt.tight_layout()
plt.show()

```

5.2 Lazy Evaluation & Memory Usage

Lazy pipeline built - nothing computed yet

Type: <class 'dask.dataframe.dask_expr._collection.Series'>

(graphviz not available - skipping task graph)


Full pipeline computed in 4.9870s

PRODUCTLINE	YEAR_ID	
Classic Cars	2003	118.363857
	2004	114.534049
	2005	112.307934
Motorcycles	2003	99.130917
	2004	98.765366
	2005	104.080392
Planes	2003	91.312561
	2004	90.395600
	2005	97.112005
Ships	2003	86.476986

Name: REVENUE_PER_UNIT, dtype: float64

Memory Usage Comparison:

Pandas (eager load)	: +-184.3 MB loaded into RAM immediately
Dask (lazy load)	: +3.4 MB (only metadata, data loaded on demand)
Memory saved	: -187.7 MB



```
[49]: #
# 5.3 - Distributed Scheduler (Multi-core)
#
print('\n 5.3 Distributed Scheduler  ')

import os
n_cores = os.cpu_count()
print(f'Available CPU cores: {n_cores}')

# Compare: synchronous vs distributed scheduler
# Operation to benchmark
def run_aggregation(ddf):
    return (
        ddf.groupby(['PRODUCTLINE', 'YEAR_ID', 'DEALSIZE'])['SALES']
        .agg(['sum', 'mean', 'count'])
        .compute()
    )

# Synchronous scheduler (single-threaded)
ddf_bench = dd.read_csv(LARGE_PATH, dtype=dtype_spec).
    ↪repartition(npartitions=best_n)

t0 = time.time()
with dask.config.set(scheduler='synchronous'):
    _ = run_aggregation(ddf_bench)
sync_time = time.time() - t0
```

```

print(f'\nSynchronous scheduler    : {sync_time:.4f}s')

# Threaded scheduler (default, uses all cores)
t0 = time.time()
with dask.config.set(scheduler='threads'):
    _ = run_aggregation(ddf_bench)
threaded_time = time.time() - t0
print(f'Threaded scheduler          : {threaded_time:.4f}s')

# Distributed scheduler (local cluster)
client = Client(
    n_workers          = min(4, n_cores),
    threads_per_worker = 2,
    memory_limit        = '2GB'
)
print(f'\nDistributed Client: {client}')
print(f'Dashboard            : {client.dashboard_link}')

t0 = time.time()
_ = run_aggregation(ddf_bench)
dist_time = time.time() - t0
print(f'Distributed scheduler    : {dist_time:.4f}s')

client.close()
print('Client closed.')

# Scheduler comparison chart
schedulers = ['Synchronous\n(single-thread)', 'Threaded\n(multi-core)',
    ↪ 'Distributed\n(local cluster)']
times      = [sync_time, threaded_time, dist_time]
colors_bar = ['firebrick', 'steelblue', 'seagreen']

fig, ax = plt.subplots(figsize=(9, 5))
bars = ax.bar(schedulers, times, color=colors_bar, width=0.5)
ax.set_title('Dask Scheduler Comparison - Synchronous vs Threaded vs ↪
    ↪ Distributed')
ax.set_ylabel('Execution Time (seconds)')
for bar, val in zip(bars, times):
    ax.text(bar.get_x() + bar.get_width()/2,
            bar.get_height() + 0.002,
            f'{val:.4f}s', ha='center', va='bottom', fontsize=11)
plt.tight_layout()
plt.show()

# Final summary
summary = pd.DataFrame({
    'Technique':      ['Partition Tuning',

```

```

        'Lazy Evaluation',
        'Memory Efficiency',
        'Distributed Scheduler'],
    'What it does': ['Finds optimal npartitions for fastest compute',
                    'Builds pipeline without loading data until .compute()',
                    'Dask loads only metadata vs Pandas loads all into RAM',
                    'Spreads work across multiple cores/workers'],
    'Result': [f'Best: {best_n} partitions ({min(partition_times):.
↵4f}s)',
               f'Pipeline computed in {elapsed:.4f}s',
               f'Saved ~{pandas_mem - dask_mem:.0f} MB vs Pandas',
               f'Distributed: {dist_time:.4f}s vs Sync: {sync_time:.
↵4f}s']
})
print('\nTask 5 - Optimization Summary:')
print(summary.to_string(index=False))

```

5.3 Distributed Scheduler

Available CPU cores: 16

Synchronous scheduler : 6.7304s

Threaded scheduler : 5.5238s

Distributed Client: <Client: 'tcp://127.0.0.1:58314' processes=4 threads=8,
memory=7.45 GiB>

Dashboard : http://127.0.0.1:8787/status


2026-05-23 17:13:24,108 - distributed.shuffle._scheduler_plugin - WARNING -
Shuffle 46c868acc096529de35bcf9b294a1d7e initialized by task ('shuffle-
transfer-46c868acc096529de35bcf9b294a1d7e', 4) executed on worker

tcp://127.0.0.1:58335

2026-05-23 17:13:24,489 - distributed.shuffle._scheduler_plugin - WARNING -
Shuffle 46c868acc096529de35bcf9b294a1d7e deactivated due to stimulus 'task-
finished-1779531204.4876149'

Distributed scheduler : 3.5136s

Client closed.



Task 5 - Optimization Summary:

Technique	What it does
Result	
Partition Tuning	Finds optimal npartitions for fastest compute
Best: 8 partitions (4.9955s)	
Lazy Evaluation	Builds pipeline without loading data until .compute()
Pipeline computed in 4.9870s	
Memory Efficiency	Dask loads only metadata vs Pandas loads all into RAM
Saved ~188 MB vs Pandas	
Distributed Scheduler	Spreads work across multiple cores/workers
Distributed: 3.5136s vs Sync: 6.7304s	

1.7 Task 6 — Reporting and Visualization Dashboard

```
[56]: import matplotlib.cm as cm

fig = plt.figure(figsize=(20, 15))
fig.suptitle('Sales Data Analytics Dashboard', fontsize=18,
             fontweight='bold', y=0.98)

# Plot 1: Monthly Sales Trend
ax1 = fig.add_subplot(3, 3, 1)
for year in sorted(monthly_sales['YEAR_ID'].unique()):
    s = monthly_sales[monthly_sales['YEAR_ID'] == year]
    ax1.plot(s['MONTH_ID'], s['SALES'], marker='o', label=str(year))
```

```

ax1.set_title('Monthly Sales Trend')
ax1.set_xlabel('Month')
ax1.set_ylabel('Sales ($)')
ax1.set_xticks(range(1, 13))
ax1.set_xticklabels(['J','F','M','A','M','J','J','A','S','O','N','D'])
ax1.legend(fontsize=8, title='Year')

# Plot 2: Sales by Product Line
ax2 = fig.add_subplot(3, 3, 2)
sales_by_product.plot(kind='bar', ax=ax2, color='steelblue')
ax2.set_title('Sales by Product Line')
ax2.set_ylabel('Sales ($)')
ax2.tick_params(axis='x', rotation=30, labelsz=8)

# Plot 3: Deal Size Distribution
ax3 = fig.add_subplot(3, 3, 3)
ax3.pie(deal_sales, labels=deal_sales.index,
        autopct='%1.1f%%', startangle=140,
        colors=['steelblue','darkorange','firebrick'])
ax3.set_title('Deal Size Distribution')

# Plot 4: Top 10 Countries
ax4 = fig.add_subplot(3, 3, 4)
sales_by_country.plot(kind='bar', ax=ax4, color='seagreen')
ax4.set_title('Top 10 Countries by Sales')
ax4.set_ylabel('Sales ($)')
ax4.tick_params(axis='x', rotation=30, labelsz=8)

# Plot 5: Customer Segmentation (fixed - uses rfm_final from Task 4)
ax5 = fig.add_subplot(3, 3, 5)
# Use k_optimal from Task 4 clustering
cluster_colors = [cm.tab10(i) for i in range(k_optimal)]
for c in range(k_optimal):
    sub = rfm_clean[rfm_clean['CLUSTER'] == c]
    ax5.scatter(sub['ORDER_COUNT'], sub['TOTAL_SALES'],
                label=cluster_labels[c],
                color=cluster_colors[c], alpha=0.6, s=30)
# Plot VIP separately
ax5.scatter(rfm_outliers['ORDER_COUNT'], rfm_outliers['TOTAL_SALES'],
            label='VIP Customers', color='gold',
            marker='*', s=150, zorder=5)
ax5.set_title('Customer Segmentation')
ax5.set_xlabel('Order Count')
ax5.set_ylabel('Total Sales ($)')
ax5.legend(fontsize=7)

# Plot 6: Quarterly Seasonality

```

```

ax6 = fig.add_subplot(3, 3, 6)
for year in sorted(quarterly_sales['YEAR_ID'].unique()):
    s = quarterly_sales[quarterly_sales['YEAR_ID'] == year]
    ax6.plot(s['QUARTER'], s['SALES'], marker='s', label=str(year))
ax6.set_title('Quarterly Seasonality')
ax6.set_xlabel('Quarter')
ax6.set_xticks([1, 2, 3, 4])
ax6.set_xticklabels(['Q1', 'Q2', 'Q3', 'Q4'])
ax6.legend(fontsize=8)

# Plot 7: Actual vs Predicted - Random Forest
ax7 = fig.add_subplot(3, 3, 7)
ax7.scatter(y_test, y_pred_rf, alpha=0.4, color='red', s=20)
ax7.plot([y_test.min(), y_test.max()],
        [y_test.min(), y_test.max()],
        'k--', linewidth=1.5, label='Perfect Prediction')
ax7.set_title('Actual vs Predicted\n(Random Forest - Order Qty)')
ax7.set_xlabel('Actual QUANTITYORDERED')
ax7.set_ylabel('Predicted QUANTITYORDERED')
ax7.legend(fontsize=7)

# Plot 8: Top 10 Customers
ax8 = fig.add_subplot(3, 3, 8)
top_customers.plot(kind='barh', ax=ax8, color='darkorange')
ax8.set_title('Top 10 Customers')
ax8.set_xlabel('Total Sales ($)')
ax8.invert_yaxis()
ax8.tick_params(axis='y', labelsize=7)

# Plot 9: Feature Importance from Random Forest (Task 4.4)
ax9 = fig.add_subplot(3, 3, 9)
importances.plot(kind='barh', ax=ax9, color='purple')
ax9.set_title('Feature Importance\n(Predict Order Quantity)')
ax9.set_xlabel('Importance Score')
ax9.tick_params(axis='y', labelsize=8)

plt.tight_layout(rect=[0, 0, 1, 0.96])
plt.savefig('output/sales_dashboard.png', dpi=150, bbox_inches='tight')
plt.show()
print('Dashboard saved → output/sales_dashboard.png')

```



Dashboard saved → output/sales_dashboard.png

1.8 Summary of Key Findings

Section	Key Finding
Data	2,823 rows × 25 columns; missing values in ADDRESSLINE2 (29.4%), STATE, POSTALCODE, and TERRITORY; 81 SALES outliers Winsorized
Top Product Line	Classic Cars generate the highest total revenue and order volume across all years
Peak Season	Q4 (October–November) consistently peaks across all years; Q1 is the lowest-performing quarter
Top Country	USA dominates total sales by a wide margin, followed by Spain and France
Top Customers	Euro Shopping Channel and Mini Gifts Distributors Ltd are the highest-spending customers

Section	Key Finding
Deal Size	Medium deals account for the largest share of orders; Large deals drive the most revenue per transaction
Customer Segments	K-Means clustering (k selected by silhouette score) identified High-Value, Mid-Value, and Low-Value segments, plus a VIP outlier group with extreme order frequency
Sales Forecast	Random Forest and Gradient Boosting models were used to predict order quantities; Random Forest achieved better R^2 with lower MAE and RMSE
Feature Engineering	Key derived features: Revenue per Unit, Discount Ratio, Quarter, Day of Week, Weekend flag — Discount Ratio and Price were top predictors
Association Rules	Apriori mining (min_support=0.05) found frequent product-line co-purchase patterns across orders
Dask Optimization	Optimal partition count determined by benchmark (lowest compute time); Dask lazy loading saved significant RAM vs Pandas eager loading; distributed scheduler tested across synchronous, threaded, and multi-worker modes

1.8.1 Business Recommendations

1. **Prioritize Classic Cars inventory and marketing** — consistently the top revenue-generating product line
2. **Launch Q3 promotions and campaigns** to drive momentum into the Q4 seasonal peak
3. **Design loyalty programs for High-Value and VIP customer segments** — they disproportionately drive total revenue
4. **Explore cross-sell bundles across frequently co-purchased product lines** based on association rule results
5. **Deepen presence in the USA, Spain, and France** — top three markets by total sales
6. **Negotiate volume pricing on Medium deals** — largest order share but lower per-transaction value than Large deals
7. **Apply discount strategically** — Discount Ratio was a significant predictor of order quantity; excessive discounting may erode margins